

→ The advantage of the method is that, at worst case scenario, convergence is exact for $k=n$. We now prove this.

⑥ Krylov subspaces and algorithm convergence

Definition: A Krylov set is defined as the set of vectors defined from an initial vector b as $\{b, Ab, A^2b, A^3b, \dots\}$

Successive Krylov subspaces are spaces spanned by successively larger groups of these vectors

$\{b\}$ ← space spanned by b

$\langle b, Ab \rangle$ " b and Ab

$\langle b, Ab, A^2b \rangle$ " "

⋮

Claims: ① The vector \underline{x}_k constructed belongs to the Krylov subspace $\langle b, Ab, \dots, A^{k-1}b \rangle$

② This subspace is equal to the subspace $\langle p_0, p_1, \dots, p_{k-1} \rangle$

Proof: By induction

for $k=1$: $\underline{x}_1 = \alpha_0 p_0 = \alpha_0 b$ ✓

Assume it's true for k : $\underline{x}_k \in \langle p_0, \dots, p_{k-1} \rangle = \langle b, Ab, \dots, A^{k-1}b \rangle$

$\underline{x}_{k+1} = \underline{x}_k + \alpha_k p_k$; so $\underline{x}_{k+1} \in \langle p_0, \dots, p_k \rangle$

$\underline{p}_k = \underline{r}_k + \beta_k p_{k-1} \Rightarrow \underline{r}_k \in \langle p_0, \dots, p_k \rangle$
 $= \underline{r}_{k-1} - \alpha_k A p_{k-1} + \beta_k p_{k-1}$

$$\begin{aligned}
 \text{So } p_k &\in \langle p_0, \dots, p_{k-1} \rangle \oplus \langle A p_{k-1} \rangle \\
 &\in \langle b, \dots, A^{k-1} b \rangle \oplus A \langle b, \dots, A^{k-1} b \rangle \\
 &\in \langle b, \dots, A^k b \rangle
 \end{aligned}$$

$$\text{So } x_{k+1} \in \langle b, \dots, A^k b \rangle \text{ too}$$

$$\text{So } \langle p_0, \dots, p_k \rangle = \langle b, \dots, A^k b \rangle.$$

Claim (3) The vector x_k minimizes $f(x)$ within the whole subspace $\langle p_0, \dots, p_{k-1} \rangle = \langle b, \dots, A^{k-1} b \rangle$

See proof in Atkinson p 565.
or in Trefethen & Bau p 296.

So this construction progressively enlarges the dimension of the space containing the guess, making sure that at every step the new guess is the global minimum within the new subspace

→ this means that addition of new dimensions to the original subspace will not change the projection of the solution on the subspaces already studied.

As a very important corollary, the last note implies that the error between x_k and the real solution x_* necessarily decreases monotonically. More precisely, it can be shown that

$$\|e_k\|_A \leq \|e_{k-1}\|_A \text{ where the norm}$$


$$\|e_k\|_A = \sqrt{(x_* - x_k)^T A (x_* - x_k)}$$

⑦ Condition number and stability issues


This section is quite general, but will also be used to discuss / introduce the idea of preconditioning.

1. Singular values

Idea: The image of the unit sphere under any $n \times n$ matrix is a hyperellipse.



$\underline{v}_1, \underline{v}_2, \underline{v}_3$
= original
unit basis



$\underline{u}_1, \underline{u}_2, \underline{u}_3$ = unit vectors along
the semi-major
axes of the hyperellipse

\underline{v}_i are determined such that

$$\boxed{A \underline{v}_i = \sigma_i \underline{u}_i}$$

$\sigma_1, \sigma_2, \sigma_3$ = "stretching factors"
= singular values
of A

Note : • If A is not a full-rank matrix,
some of the singular values are 0.

(if A is rank r , exactly r singular values
are non-zero).

• By construction $\sigma_i \geq 0$

• By convention, the \underline{u}_i and \underline{v}_i are
usually ordered such that
 $\sigma_1 \geq \sigma_2 \geq \sigma_3 \dots \geq \sigma_n \geq 0$

• This concept is entirely related to the
Singular Value Decomposition (see Tidtem
project).

2 First notion of ill-conditioning

when applying a matrix A to a vector x , imagine first writing \underline{x} on the basis of the vectors \underline{v}_k

$$\underline{x} = \sum \alpha_k \underline{v}_k$$

$$\text{so } A\underline{x} = \sum_k \alpha_k \sigma_k \underline{u}_k$$

Even if A is not singular, problems may arise if $\sigma_1 \gg \sigma_n$. Indeed, in that case the term $\alpha_n \sigma_n \underline{u}_n$ may be negligible in front of $\alpha_1 \sigma_1 \underline{u}_1$ and roundoff errors may affect it.

Geometrically speaking, it is particularly easy to visualize when $\sigma_n \ll \sigma_{n-1} \sim \sigma_{n-2} \dots \sigma_1$
(the hyperellipse is very flat in one direction)

$$\text{or } \sigma_1 \gg \sigma_2 \sim \sigma_3 \dots \sigma_n$$

(the hyperellipse is very elongated in one direction)

\Rightarrow in both cases (and any intermediate case $\sigma_1 \sim \sigma_2 \sim \dots \sigma_k \gg \sigma_{k+1} \sim \dots \sigma_n$)

information is lost if roundoff errors occur.

\rightarrow we expect the sensitivity of a problem to small roundoff errors to be particularly dependent on the value of

$$\kappa = \frac{\sigma_1}{\sigma_n}$$

We now formalize this idea in more mathematical terms.

3. Necessary mathematical tool: the norm of a matrix

We saw that the norm of a vector can be defined in many ways, provided it satisfies

- $\|x\| \geq 0 \quad \forall x$ and $\|x\| = 0 \Leftrightarrow x = 0$
- $\|x+y\| \leq \|x\| + \|y\|$
- $\|\alpha x\| = |\alpha| \|x\|$.

A particularly useful norm is the Euclidean norm:

$$\begin{aligned}\|x\|_2 &= \sqrt{x^T x} \\ &= \left(\sum_{i=1}^n |x_i|^2 \right)^{1/2}\end{aligned}$$

(see handout for other norms)

We now define the norm of a matrix as the maximum factor by which a matrix A can stretch a vector x

$$\|A\| = \sup_x \frac{\|Ax\|}{\|x\|} = \sup_{\|x\|=1} \|Ax\|.$$

(note that this holds even when A is not square)

Now by using the Euclidean norm and writing $x = \sum_k \alpha_k v_k$

it is easy to show that $\boxed{\|A\|_2 = \sigma_1}$

In addition, one can also show that

$$\|A^{-1}\|_2 = \frac{1}{\sigma_n}$$

4. Mathematical definition of conditioning

① Idea: Given a mathematical function

$$\begin{aligned} f: \mathbb{R}^m &\rightarrow \mathbb{R}^n \\ \underline{x} &\rightarrow f(\underline{x}) \end{aligned}$$

how do perturbations in the input \underline{x} affect the output $f(\underline{x})$?

- If small "errors" in \underline{x} result only in equivalently small "errors" in $f(\underline{x})$ then the problem is well-conditioned.
- If small "errors" in \underline{x} result in much larger "errors" in $f(\underline{x})$ then the problem is ill-conditioned.

let's define the relative condition number as

$$\kappa(\underline{x}) = \lim_{\delta \rightarrow 0} \sup_{\|\delta \underline{x}\| \leq \delta} \frac{\|\delta f\| / \|f\|}{\|\delta \underline{x}\| / \|\underline{x}\|} \quad \leftarrow \frac{\text{relative change in } f}{\text{rel. change in } \underline{x}}$$

$$\text{where } \delta f = f(\underline{x} + \delta \underline{x}) - f(\underline{x})$$

$\Rightarrow \kappa(\underline{x})$ large (for some \underline{x}) means that (for that \underline{x}), a small change $\delta \underline{x}$ results in a large change $\|\delta f\|$.

If $f(\underline{x})$ is a differentiable function, then

$$\delta f = J(\underline{x}) \delta \underline{x}$$

\uparrow Jacobian matrix.

$$\text{So } \sup_{\|\delta \underline{x}\| \leq \delta} \frac{\|\delta f\|}{\|\delta \underline{x}\|} = \sup_{\|\delta \underline{x}\| \leq \delta} \frac{\|J(\underline{x}) \delta \underline{x}\|}{\|\delta \underline{x}\|} = \|J(\underline{x})\| \quad \text{by definition}$$

$$\text{and } \kappa(\underline{x}) = \|J(\underline{x})\| \cdot \frac{\|\underline{x}\|}{\|f(\underline{x})\|}$$

② Conditionning of the problem : what is the error on Ax when A is fixed?

Ax is a linear function so $J \equiv A$

$$\begin{aligned}\Rightarrow \kappa(x) &= \|J(x)\| \cdot \frac{\|x\|}{\|Ax\|} \\ &= \|A\| \cdot \frac{\|x\|}{\|Ax\|}\end{aligned}$$

If we use the 2-norm, $\|x\|_2 = \|A^{-1}Ax\|_2$
(by Cauchy's inequality) $\leq \|A^{-1}\| \|Ax\|_2$

$$\text{so } \kappa_2(x) \leq \|A\|_2 \|A^{-1}\|_2 = \frac{\sigma_1}{\sigma_m}$$

and, for some x the upper bound is actually achieved.

\rightarrow this confirms our expectations that $\frac{\sigma_1}{\sigma_m}$ is a good descriptor of the conditionning of a problem.

The quantity $\|A\| \|A^{-1}\| \equiv \kappa(A)$ is called the condition number of A .

If A is singular, by convention $\kappa(A) = \infty$.

③ Conditionning of a system of equations w.r.t. perturbations in A .

Let's now suppose we are designing an algorithm working on A which progressively introduces small roundoff errors. How is $x = A^{-1}b$ affected?

\rightarrow Similarly, it is possible to show that the condition number is $\kappa(A) = \|A\| \|A^{-1}\|$.

What this implies in practise is that an algorithm working on A introducing relative error $\frac{\|SA\|}{\|A\|}$ no larger than $O(\epsilon_{\text{machine}})$ in the coefficients result in an error on the outcome

$$\frac{\|Sx\|}{\|x\|} = O(\kappa(A) \epsilon_{\text{machine}}).$$

→ if $\kappa(A)$ is very large, many significant digits in the solution are lost

Note that this assumes roundoff errors themselves are only $O(\epsilon_{\text{machine}})$. Some algorithms (cf G-E without pivoting) introduce roundoff errors $\gg \epsilon_{\text{machine}}$!

5. In practise, how to evaluate $\kappa(A)$?

- Calculating A^{-1} , then $\|A^{-1}\|$ is too CPU-expensive
- See LAPACK routines `**CON` (specialized) or `DGESVX.f` (driver routine) for an estimate of the reciprocal of the condition number (returned in RCOND)
- Note that the Lapack routines return either the reciprocal of $\kappa_1(A) = \|A\|_1 \|A^{-1}\|_1$, or $\kappa_\infty(A) = \|A\|_\infty \|A^{-1}\|_\infty$

where $\|A\|_1$ is the norm of A based on $\|x\|_1$,
 $(\|x\|_1 = \sum_i |x_i|)$

and $\|A\|_\infty$ is the norm of A based on $\|x\|_\infty$
 $(\|x\|_\infty = \max_i |x_i|)$

Example

$$A = \begin{pmatrix} 1+10^{-k} & 1 \\ 1 & 1 \end{pmatrix}$$

$$\|A\|_1 = \max_j \sum_i |a_{ij}| = \max(2+10^{-k}, 2) = 2+10^{-k}$$

$$\|A\|_\infty = \max_i \sum_j |a_{ij}| = 2+10^{-k} \quad (\text{same})$$

look for A^{-1} :

$$\left(\begin{array}{cc|cc} 1+10^{-k} & 1 & 1 & 0 \\ 1 & 1 & 0 & 1 \end{array} \right)$$

$$\rightarrow \left(\begin{array}{cccc} 1 & \frac{1}{1+10^{-k}} & \frac{1}{1+10^{-k}} & 0 \\ 0 & \frac{10^{-k}}{1+10^{-k}} & -\frac{1}{1+10^{-k}} & 1 \end{array} \right)$$

$$\rightarrow \left(\begin{array}{cccc} 1 & 0 & \frac{1+10^k}{1+10^{-k}} & -10^k \\ 0 & 1 & -\frac{1}{10^{-k}} & \frac{1+10^{-k}}{10^{-k}} \end{array} \right)$$

$$\text{so } A^{-1} = \begin{pmatrix} \frac{1+10^k}{1+10^{-k}} & -10^k \\ -10^k & 1+10^k \end{pmatrix}$$

$$\text{so } \|A^{-1}\|_1 \approx 2+10^k$$
$$\|A^{-1}\|_\infty \approx 2+10^k$$

$$\Rightarrow \boxed{\kappa(A) \approx 2 \cdot 10^k \text{ for large } k}$$

Note that if $b = \begin{pmatrix} 1 \\ 1 \end{pmatrix}$ then

$$x = A^{-1}(b) = \begin{pmatrix} 0 \\ 1 \end{pmatrix}$$

$$\text{if } b = \begin{pmatrix} 1+\epsilon \\ 1 \end{pmatrix} \text{ then } A^{-1}b \approx \begin{pmatrix} \epsilon 10^k \\ 1 \end{pmatrix}$$

③ Convergence of conjugate Gradient algorithm

In fact, it can be shown that

$$\|e_n\|_A \leq 2 \left(\frac{\sqrt{\kappa} - 1}{\sqrt{\kappa} + 1} \right)^n \|e_0\|_A \quad \left(\text{see Trefethen \& Bau} \right) \cdot$$

p 300

$$\text{where } \kappa = \frac{\lambda_{\max}}{\lambda_{\min}} = \frac{\text{largest eigenvalue}}{\text{smallest eigenvalue}}.$$

(note that by assumption A is positive definite so all the eigenvalues are > 0)

- For quick convergence, we would hope that κ is "not too large" (ideally we want $\frac{\sqrt{\kappa} - 1}{\sqrt{\kappa} + 1} \ll 1$)

in practise typical $\kappa \cong 1 \rightarrow 10$ is great!

- For ill-conditioned matrices, the gain of iterative methods vs standard direct methods is negligible

\Rightarrow IDEA OF PRE-CONDITIONING

⑨ Preconditioning ideas

- The convergence rate of an iterative method depends on the condition number $\kappa = \frac{\lambda_{\max}}{\lambda_{\min}}$ of A .
- For ill-conditioned matrices $\kappa \gg 1$ and the convergence can be v. slow.
- Idea: Instead of solving $Ax=b$, solve the equivalent system

$$KAX = Kb \quad \text{where}$$

K = a non singular matrix

KA = a better conditioned matrix

- Note: Although the product KA is never actually formed in a preconditioned algorithm, it is vital that the product Kv or $K^T v$ or $K^{-1}v$ be fast for any vector v ($a < n^2$ process)
- ↳ K should be a simple or v. sparse matrix.

- Finding preconditioners that satisfy both requirements is ongoing research

- different types of matrices A respond better to different types of preconditioner

- Preconditioning can be mathematically based or physically based

Mathematical

Diagonal
Incomplete Cholesky (Golub & V. Loan)

Physical

Coarse grid/Multigrid
low-pass filters

See Appendix on preconditioning for actual implementation

⑩ Generalization to non-symmetric matrices

The idea of minimizing $f(\underline{x}_k)$ for $\underline{x}_k \in S_k$
(the Krylov subspace)

$$S_k = \langle b, Ab, \dots, A^{k-1}b \rangle$$

for successively larger values of k can be
generalized to non-symmetric matrices.

The methods are known as GMRES
(Generalized Minimum Residuals).

See Trefethen & Bau Lecture 35.